

Description

You are asked to make predictions for each unique id in the test dataset about the likelihood of the person having a bank account or not, i.e. Yes = 1, No = 0. You will train your model on 70% of the data and test your model on the final 30% of the data.

In [0]:

```
# Importing necessary Libraries to use
import pandas as pd
import numpy as np
import seaborn as sns
sns.set()
import matplotlib.pyplot as plt
```

In [0]:

```
train = pd.read_csv("Train_v2.csv")
test = pd.read_csv ('Test_v2.csv')
```

In [0]:

```
#DATA CLEANING ON TRAIN DATASET
```

In [264]:

```
train.head()
```

Out[264]:

	country	year	unique_id	bank_account	location_type	cellphone_access	household_size	age_of_respondent	gender_of_respondent	relationship_with_head	marital_status	education_level	job_type
0	Kenya	2018	unique_id_1	Yes	Rural	Yes	3	24	Female	Spouse	Married/Living together	Secondary education	Self employed
1	Kenya	2018	unique_id_2	No	Rural	No	5	70	Female	Head of Household	Widowed	No formal education	Government Dependent
2	Kenya	2018	unique_id_3	Yes	Urban	Yes	5	26	Male	Other relative	Single/Never Married	Vocational/Specialised training	Self employed
3	Kenya	2018	unique_id_4	No	Rural	Yes	5	34	Female	Head of Household	Married/Living together	Primary education	Formally employed

co un tr y	y e a r	uni que id	bank _acc ount	locat ion_ type	cellph one_a ccess	hous ehold _size	age_of _respo ndent	gender_ of_resp ondent	relations hip_with _head	mari tal_s tatus	educati on_leve l	job _ty pe
---------------------	------------------	------------------	----------------------	-----------------------	--------------------------	------------------------	---------------------------	------------------------------	--------------------------------	------------------------	-------------------------	------------------

d
Priv
ate

4	Kenya	2018	uniqueid_5	No	Urban	No	8	26	Male	Child	Single/Never Married	Primary education	Informally employed
---	-------	------	------------	----	-------	----	---	----	------	-------	----------------------	-------------------	---------------------

In [265]:

```
train.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23524 entries, 0 to 23523
Data columns (total 13 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   country                             23524 non-null  object
 1   year                                23524 non-null  int64
 2   uniqueid                            23524 non-null  object
 3   bank_account                        23524 non-null  object
 4   location_type                       23524 non-null  object
 5   cellphone_access                    23524 non-null  object
 6   household_size                      23524 non-null  int64
 7   age_of_respondent                   23524 non-null  int64
 8   gender_of_respondent                 23524 non-null  object
 9   relationship_with_head               23524 non-null  object
10   marital_status                       23524 non-null  object
11   education_level                     23524 non-null  object
12   job_type                            23524 non-null  object
dtypes: int64(3), object(10)
memory usage: 2.3+ MB
```

In [266]:

```
# Checking the null values from the train dataset
train.isna().sum()
```

Out[266]:

country	0
year	0
uniqueid	0
bank_account	0
location_type	0
cellphone_access	0
household_size	0

```
age_of_respondent      0
gender_of_respondent    0
relationship_with_head  0
marital_status          0
education_level         0
job_type                0
dtype: int64
```

In [267]:

```
# now it time to remove outliers in our train dataset - we are using z-score to detect and remove the outliers
```

```
from scipy import stats
```

```
z = np.abs(stats.zscore(train._get_numeric_data()))
```

```
z
```

Out[267]:

```
array([[1.20854126, 0.35800673, 0.89618796],
       [1.20854126, 0.53983446, 1.88827897],
       [1.20854126, 0.53983446, 0.77512418],
       ...,
       [1.20854126, 0.53983446, 0.71459229],
       [1.20854126, 1.43767565, 0.53299662],
       [1.20854126, 2.78443744, 1.13831551]])
```

In [0]:

```
train=train[(z<3).all(axis=1)]
```

In [269]:

```
train.shape
```

Out[269]:

```
(23232, 13)
```

In [270]:

```
train.describe() # 3 numerical columns only. Most of the columns are categorical.
```

Out[270]:

	year	household_size	age_of_respondent
count	23232.000000	23232.000000	23232.000000
mean	2016.969697	3.733815	38.610064
std	0.846098	2.095128	16.173751
min	2016.000000	1.000000	16.000000
25%	2016.000000	2.000000	26.000000

	year	household_size	age_of_respondent
50%	2017.000000	3.000000	35.000000
75%	2018.000000	5.000000	49.000000
max	2018.000000	10.000000	88.000000

In [271]:

```
train.dtypes
```

Out[271]:

```
country          object
year             int64
uniqueid         object
bank_account     object
location_type    object
cellphone_access object
household_size   int64
age_of_respondent int64
gender_of_respondent object
relationship_with_head object
marital_status   object
education_level  object
job_type         object
dtype: object
```

In [272]:

```
# Lets see if our dataset is balanced or not by checking our target distribution
train.bank_account.value_counts()
```

Out[272]:

```
No      19946
Yes      3286
Name: bank_account, dtype: int64
```

In [273]:

```
a = len(train[train.bank_account=='Yes'])
b = len(train[train.bank_account=='No'])
c = len(train)
print('We have an imbalanced dataset with a %i/%i ratio'%((b/c*100),(a/c*100)+1))
```

We have an imbalanced dataset with a 85/15 ratio
The id+country thing

In [274]:

```
train.uniqueid.value_counts().head(5) #
```

Out[274]:

```
uniqueid_1502      4
```

```
uniqueid_985      4
uniqueid_18       4
uniqueid_1746     4
uniqueid_12       4
Name: uniqueid, dtype: int64
```

In [0]:

```
# some uniqueid's have the same value 4 times. this can be explained by the
number of countries in the dataset. Hence the need to
# identify people by 'country'+ 'uniqueid' to avoid having duplicate uniqueid
values.
```

In [276]:

```
test.uniqueid.value_counts().head(5)
```

Out[276]:

```
uniqueid_8612      3
uniqueid_8627      3
uniqueid_8650      3
uniqueid_8633      3
uniqueid_8605      3
Name: uniqueid, dtype: int64
```


In [0]:

```
# Same thing needs to be done for the test set.
```

Exploratory Data Analysis

In [278]:

```
# age_of_respondant
hist_age = train.age_of_respondent.hist(bins=25,figsize=[15,10])
```



In [0]:

We have a skewed to the right distribution for this variable.


Country

In [280]:

```
train['country'].value_counts().plot(kind='bar', figsize=[15,5])
```

Out[280]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d7e20668>



In [0]:

Rwanda is the most occurring value for country, while Uganda is the least occurring one.


location type

In [282]:

```
train['location_type'].value_counts().plot(kind='bar', figsize=[15,5])  
#plot()
```

Out[282]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d7dfd7b8>



In [0]:

More people from rural places have been interviewed than people in urban places.


cellphone access

In [284]:

```
train['cellphone_access'].value_counts().plot(kind='bar', figsize=[15,5])
```

Out[284]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d7d4c8d0>




In [0]:

'Yes' indicates the non-possession of a bank account which is unlikely.

household_size

In [286]:

```
hist_hs = train.household_size.hist(bins=20,figsize=[15,10])
```



In [0]:

Another numerical distribution that's skewed to the right


gender of respondent

In [288]:

```
train['gender_of_respondent'].value_counts().plot(kind='bar',figsize=[15,5])
```

Out[288]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d848ba90>



In [0]:

```
# Trainset has more females than males.
```


relationship with head

In [290]:

```
train['relationship_with_head'].value_counts().plot(kind='bar',figsize=[15,5])
```

Out[290]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d848b4e0>



In [0]:

```
# 6 category and the most occuring is 'Head of Household' followed by Spouse.
```


marital status

In [292]:

```
train['marital_status'].value_counts().plot(kind='bar',figsize=[15,5])
```

Out[292]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95dad8aa20>



In [0]:

5 categories with one category 'Don't know' being significantly undersampled


education_level

In [294]:

```
train['education_level'].value_counts().plot(kind='bar', figsize=[15,5])
```

Out[294]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95dac035c0>



In [0]:

7 categories, one category is called '6' and is undersampled; another category is 'Other/Dont know/RTA' is also undersampled


Job_type

In [296]:

```
train['job_type'].value_counts().plot(kind='bar', figsize=[15,5])
```

Out[296]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d7c2eda0>



Let's check the year variable

In [297]:

```
test.year.value_counts()
```

Out[297]:

```
2016    3745
2018    3502
2017    2839
Name: year, dtype: int64
```

In [298]:

```
train.year.value_counts()
```

Out[298]:

```
2016    8678
2018    7974
2017    6580
Name: year, dtype: int64
```

In [0]:

```
# Same thing, the year seems to have been a condition to be respected when
splitting the train/test
```

In [300]:

```
train[train.year==2016].country.value_counts()
```

Out[300]:

```
Rwanda    8678
Name: country, dtype: int64
```

In [301]:

```
train[train.year==2017].country.value_counts()
```

Out[301]:

```
Tanzania    6580
Name: country, dtype: int64
```

In [302]:

```
train[train.year==2018].country.value_counts()
```

Out[302]:

```
Kenya      5978
Uganda     1996
Name: country, dtype: int64
```

In [0]:

```
# The year variable is indicative of which country is mentionned. That's how the train set is made.
```

In [0]:

```
train['bank_account'].replace({'No': 0, 'Yes': 1}, inplace = True)
```

In [0]:

```
#labels = ['0-19', '20-29', '30-39', '40-49', '50-59', '60-69', '70-79', '80-100']
#train['age_group'] = pd.cut(train.age_of_respondent, range(0,81,10), right=False, labels=labels)
```

age_group

In [0]:


```
#plt.figure(figsize=[18,12])
#sns.barplot('age_group', 'bank_account', data=train)
```

In [307]:

```
plt.figure(figsize=[15,6])
sns.barplot('gender_of_respondent', 'bank_account', data=train)
```

Out[307]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d7c2ec50>



In [0]:

```
# Males are more likely to have a bank account according to this plot
```

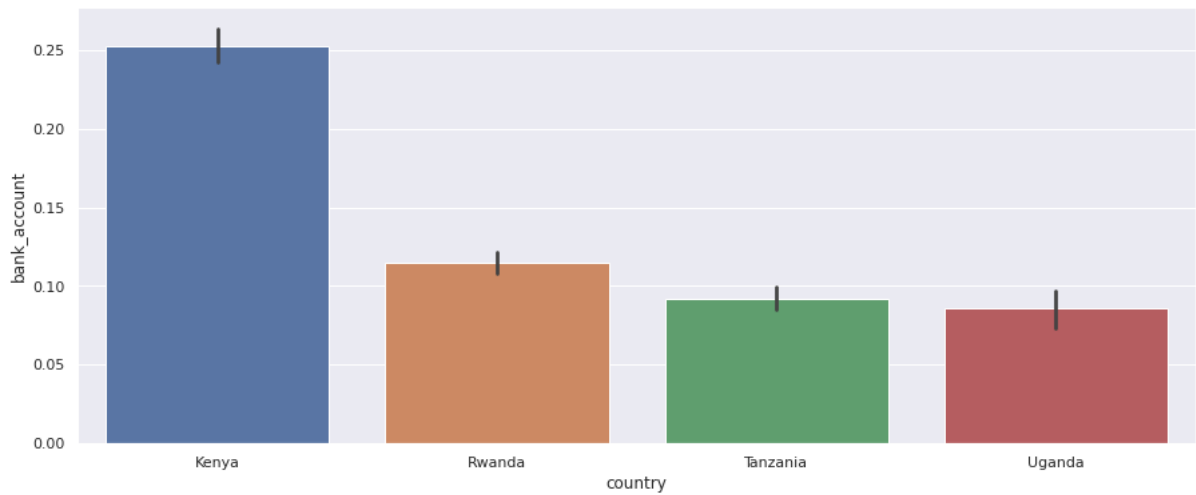
Country

In [309]:

```
plt.figure(figsize=[15,6])
sns.barplot('country', 'bank_account', data=train)
```

Out[309]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d7c256a0>




Job Type

In [310]:

```
plt.figure(figsize=[25,6])
sns.barplot('job_type', 'bank_account', data=train)
```

Out[310]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d7af7d68>



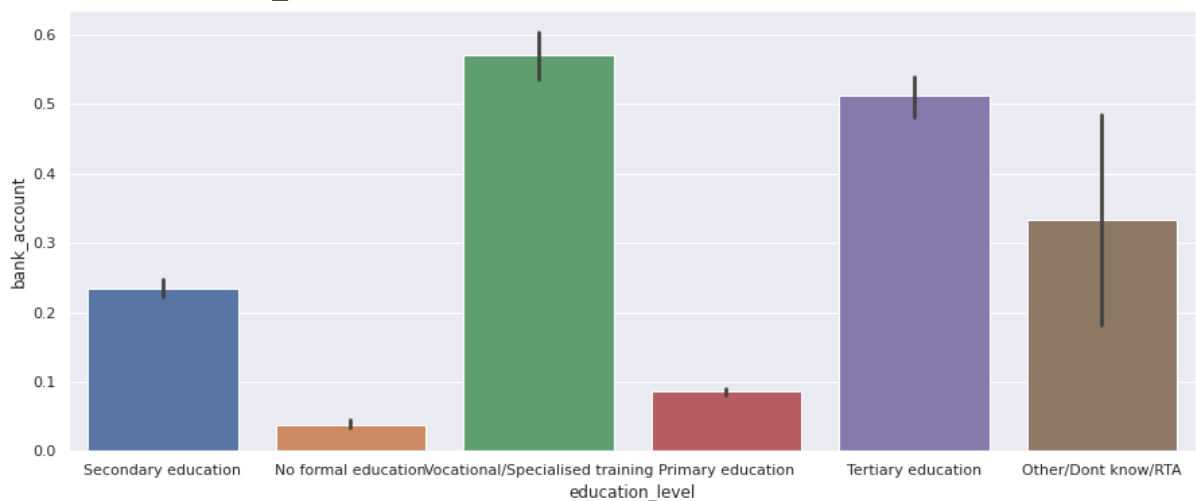
Education_Level

In [311]:

```
plt.figure(figsize=[15,6])
sns.barplot('education_level', 'bank_account', data=train)
```

Out[311]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d7a190b8>




Household_Size

In [312]:

```
plt.figure(figsize=[15,6])
sns.barplot('household_size', 'bank_account', data=train)
```

Out[312]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d7af7b38>




Relationship_with_head

In [313]:

```
plt.figure(figsize=[15,6])
sns.barplot('relationship_with_head', 'bank_account', data=train)
```

Out[313]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d791d390>



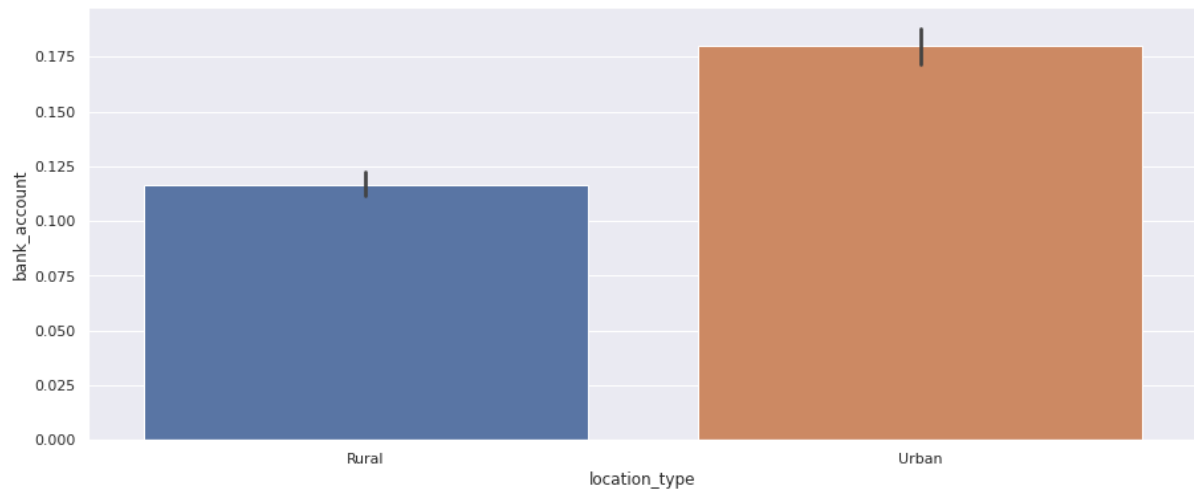
Location_Type

In [314]:

```
plt.figure(figsize=[15,6])
sns.barplot('location_type', 'bank_account', data=train)
```

Out[314]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d7917860>



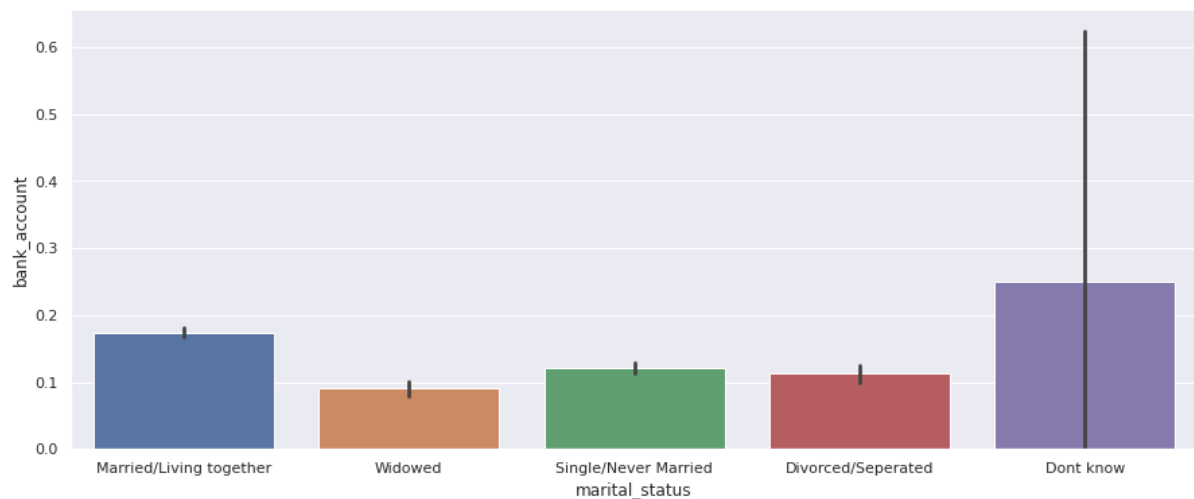
Marital Status

In [315]:

```
plt.figure(figsize=[15,6])
sns.barplot('marital_status', 'bank_account', data=train)
```

Out[315]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f95d798a5c0>



In [0]:

DATA MODELING

In [0]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn import model_selection, preprocessing

from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn import metrics
```

In [0]:

```
#Data Conversion for the Train Dataset
for e in train.columns:
    if train[e].dtype == 'object':
        lbl = preprocessing.LabelEncoder()
        lbl.fit(list(train[e].values))
        train[e] = lbl.transform(list(train[e].values))
```

In [318]:

```
train.head()
```

Out[318]:

	co un try	y e a r	uni qu eid	bank _acco unt	locat ion_t ype	cellph one_a ccess	house hold_ size	age_of _respo ndent	gender_ of_respo ndent	relations hip_with _head	marit al_st atus	educa tion_l evel	job _ty pe
0	0	2018	0	1	0	1	3	24	0	5	2	3	9
1	0	2018	1111	0	0	0	5	70	0	1	4	0	4
2	0	2018	2222	1	1	1	5	26	1	3	3	5	9
3	0	2018	3333	0	0	1	5	34	0	1	2	2	3
4	0	2018	4444	0	1	0	8	26	1	0	3	2	5

In [319]:

```
train['country'].unique()
```

Out[319]:

```
array([0, 1, 2, 3])
```

In [0]:

```
train = train.drop(columns=['year', 'uniqueid', 'household_size', 'education_level', 'marital_status', 'relationship_with_head'], axis=1)
```

In [321]:

```
train.head()
```

Out[321]:

	country	bank_account	location_type	cellphone_access	age_of_respondent	gender_of_respondent	job_type
0	0	1	0	1	24	0	9
1	0	0	0	0	70	0	4
2	0	1	1	1	26	1	9
3	0	0	0	1	34	0	3
4	0	0	1	0	26	1	5

Train and test split

In [0]:

```
#X_train, X_test, y_train, y_test = train_test_split(X,Y , test_size = 0.3)
# Splitt+ing the data from the train dataset
```

In [0]:

```
X = train.loc[:,train.columns!='bank_account']
```

In [0]:

```
Y = train['bank_account']
```

In [0]:

```
from sklearn.model_selection import train_test_split
```

In [0]:

```
X_train, X_test, Y_train, Y_test = train_test_split(X,Y , test_size = 0.3)
```

In [327]:

```
X_train.shape, X_test.shape
```

Out[327]:

```
((16262, 6), (6970, 6))
```

Build the model on training data

In [0]:

```
# Hypothesis statement
```

```
# Ho is not significant
```

```
# HR is significant
```

```
# P-value < 5%
```

In [0]:

```
# import statsmodels.api as sm
```

In [0]:

```
# model_1 = sm.OLS(Y_train,X_train).fit()
```

In [0]:

```
# model_1.summary2()
```


In [0]:

```
# P-value < 5%
# HR: True
# H0: False
```

Logistic Regression Model

In [333]:

```
#Training
from sklearn.linear_model import LogisticRegression
modell = LogisticRegression()
modell.fit(X_train,Y_train)
```

Out[333]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True
,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0
,
                    warm_start=False)
```

In [334]:

```
#Testing
predicted = modell.predict(X_test)
predicted
```

Out[334]:

```
array([0, 0, 0, ..., 0, 0, 0])
```

In [335]:

```
#Evaluation
#Confusion Matrix
print(metrics.confusion_matrix(Y_test, predicted))
[[5956  16]
 [ 978  20]]
```

In [336]:

```
#Classification Report
print("\nAcuracy Score of LogisticRegression Model:")
print(metrics.accuracy_score(Y_test, predicted))
print("\nClassification Report:")
print(metrics.classification_report(Y_test, predicted))
```

```
Acuracy Score of LogisticRegression Model:
0.8573888091822095
```

Classification Report:

	precision	recall	f1-score	support
0	0.86	1.00	0.92	5972
1	0.56	0.02	0.04	998
accuracy			0.86	6970

macro avg	0.71	0.51	0.48	6970
weighted avg	0.82	0.86	0.80	6970


In [337]:

```
# Using the ROC Curve to See the accuracy of our Model
print("ROC Curve")
modell_prob = modell.predict_proba(X_test)
modell_prob1 = modell_prob[:,1]
fpr,tpr,thresh = metrics.roc_curve(Y_test,modell_prob1)
roc_auc_lr = metrics.auc(fpr,tpr)
plt.figure(dpi=80)
plt.title("ROC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.plot(fpr,tpr, 'r',label = 'AUC Score = %0.2f'%roc_auc_lr)
plt.plot(fpr,fpr, 'b--',color='blue')
plt.legend()
```

ROC Curve

Out[337]:

<matplotlib.legend.Legend at 0x7f95d77182b0>



Random Forest Model

In [338]:

```
#Training
from sklearn.ensemble import RandomForestClassifier
model2 = RandomForestClassifier()
model2.fit(X_train,Y_train)
```

Out[338]:

```

RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                       criterion='gini', max_depth=None, max_features='auto',
                       max_leaf_nodes=None, max_samples=None,
                       min_impurity_decrease=0.0, min_impurity_split=None,
                       min_samples_leaf=1, min_samples_split=2,
                       min_weight_fraction_leaf=0.0, n_estimators=100,
                       n_jobs=None, oob_score=False, random_state=None,
                       verbose=0, warm_start=False)

```

In [339]:

```

#Testing
predicted = model2.predict(X_test)
predicted

```

Out[339]:

```

array([0, 0, 0, ..., 0, 0, 0])

```

In [340]:

```

#Evaluation
#Confusion Matrix
print(metrics.confusion_matrix(Y_test, predicted))

[[5616  356]
 [ 662  336]]

```

In [341]:

```

#Classification Report
print("\nAcuracy Score of RF Model:")
print(metrics.accuracy_score(Y_test, predicted))
print("\nClassification Report:")
print(metrics.classification_report(Y_test, predicted))

Acuracy Score of RF Model:
0.853945480631277

```

```

Classification Report:

```

	precision	recall	f1-score	support
0	0.89	0.94	0.92	5972
1	0.49	0.34	0.40	998
accuracy			0.85	6970
macro avg	0.69	0.64	0.66	6970
weighted avg	0.84	0.85	0.84	6970

In [342]:

```

print("ROC Curve")
model2_prob = model2.predict_proba(X_test)
model2_prob1 = model2_prob[:,1]
fpr,tpr,thresh = metrics.roc_curve(Y_test,model2_prob1)
roc_auc_rf = metrics.auc(fpr,tpr)


```

```
plt.figure(dpi=80)
plt.title("ROC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.plot(fpr, tpr, 'r', label = 'AUC Score = %0.2f'%roc_auc_rf)
plt.plot(fpr, fpr, 'b--', color='blue')
plt.legend()
```

ROC Curve

Out[342]:

<matplotlib.legend.Legend at 0x7f95d768c908>



KNN Model

In [343]:

```
#Training
from sklearn import neighbors
model3 = neighbors.KNeighborsClassifier()

model3.fit(X_train,Y_train)
```

Out[343]:

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                    weights='uniform')
```

In [344]:

```
#Testing
predicted = model3.predict(X_test)
predicted
```

Out[344]:

```
array([0, 0, 0, ..., 0, 0, 0])
```

In [345]:

```
#Evaluation
#Confusion Matrix
print(metrics.confusion_matrix(Y_test, predicted))

[[5741  231]
 [ 704  294]]
```

In [346]:

```
#Classification Report
print("\nAccuracy Score of KNN Model:")
print(metrics.accuracy_score(Y_test, predicted))
print("\nClassification Report:")
print(metrics.classification_report(Y_test, predicted))

Accuracy Score of KNN Model:
0.8658536585365854
```

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.96	0.92	5972
1	0.56	0.29	0.39	998
accuracy			0.87	6970
macro avg	0.73	0.63	0.66	6970
weighted avg	0.84	0.87	0.85	6970

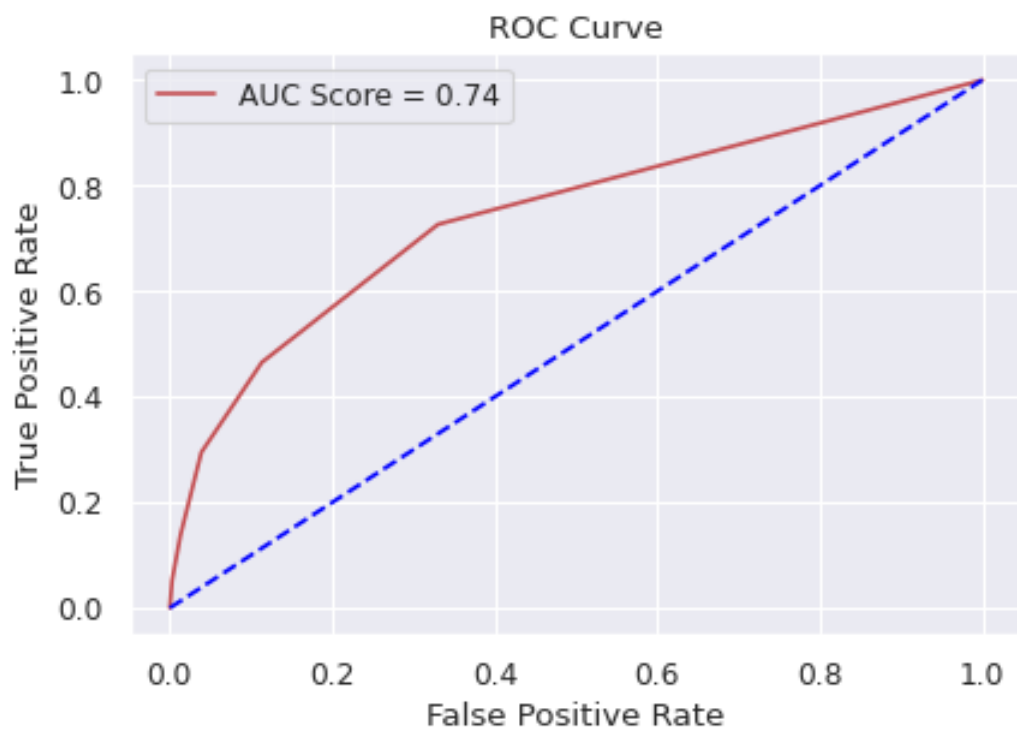
In [347]:

```
print("ROC Curve")
model3_prob = model3.predict_proba(X_test)
model3_prob1 = model3_prob[:,1]
fpr,tpr,thresh = metrics.roc_curve(Y_test,model3_prob1)
roc_auc_knn = metrics.auc(fpr,tpr)
plt.figure(dpi=80)
plt.title("ROC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.plot(fpr,tpr,'r',label = 'AUC Score = %0.2f'%roc_auc_knn)
plt.plot(fpr,fpr,'b--',color='blue')
plt.legend()

ROC Curve
```

Out[347]:

<matplotlib.legend.Legend at 0x7f95d76746a0>



In [0]:

In [0]: